

Device Management

Assignment Project Exam Help

(w <https://eduassistpro.github.io>
axgopala@imperial.ac.uk

Add WeChat edu_assist_pr

Assignment Project Exam Help

Fair access to shared devices

- Allocation of dedicated devices

Expl

<https://eduassistpro.github.io>

Provide uniform simple view of I/O

- Hide complexity of device handling
- Give uniform naming and error handling

Add WeChat edu_assist_pr

Block Devices

- Stores information in fixed-size blocks
- Transfers are in units of entire

ch <https://eduassistpro.github.io>

- Delivers or accepts stream of characters, without regard to block structure
- Not addressable, does not have any seek operation

How does the OS actually communicate with the hardware?

Each hardware controller has a few registers used for communication with the CPU

Assignment Project Exam Help

OS can write to these registers to command the device

- Deliver data



• <https://eduassistpro.github.io>

- Perform some action

OS can read from these registers to learn about

- State of the device
- Whether it is ready to accept commands
- ...

Device independence from

- Device type (e.g. terminal, disk or DVD drive)
- Device instance (e.g. which disk)

Uniform naming name of a file should be a string or integer and n

Devi

- Unit of data transfer: **character** or **block**
- Supported operations: e.g. **read**, **write**, **see**
- **Synchronous** or **asynchronous** operation

- Speed differences
- Sharable (e.g. disks) or single user (e.g. printer, DVD-RW)
- Error handling
- Buffering

Assignment Project Exam Help

Programmed I/O

Inte

Dire

<https://eduassistpro.github.io>

Example used: Printing a string

Add WeChat edu_assist_pro

Assignment Project Exam Help

<https://eduassistpro.github.io>

```
copy_from_user (buffer, p, count); // p = kernel buffer
for (i = 0; i < count; i++) {      // loop on every char
    while (*printer_status_reg != READY); // loop until ready
    *printer_data_register = p[i];      // output one char
}
return_to_user ();
```

Assignment Project Exam Help

<https://eduassistpro.github.io>

Add WeChat edu_assist_pr

The connections between the devices and the interrupt controller actually use interrupt lines on the bus rather than dedicated wires


```
copy_from_user (buffer, p, count);  
enable_interrupts ();  
while (*printer_status_reg != READY);  
*printer_data_register = p[0];  
scheduler ();
```

Assignment Project Exam Help

<https://eduassistpro.github.io>

```
} else {  
    *printer_data_register = p[i];  
    count = count - 1;  
    i++;  
}  
acknowledge_interrupt ();  
return_from_interrupt ();
```

Add WeChat edu_assist_pro

Interrupt service procedure for the printer

Assignment Project Exam Help

<https://eduassistpro.github.io>

Add WeChat edu_assist_pr

Operation of a DMA transfer

Assignment Project Exam Help

```
copy_from_user (buffer, p, count);  
set_up_DMA_controller ();  
scheduler ();
```

<https://eduassistpro.github.io>

```
acknowledge_interrupt ();  
unblock_user ();  
return from_interrupt ();
```

Add WeChat edu_assist_pro

Interrupt service procedure

Assignment Project Exam Help

<https://eduassistpro.github.io>

Add WeChat edu_assist_pr

Assignment Project Exam Help

<https://eduassistpro.github.io>

Add WeChat edu_assist_pr

Assignment Project Exam Help

<https://eduassistpro.github.io>

Add WeChat edu_assist_pr

- 1 Driver starts an I/O operation block and blocks until the I/O has completed

2

<https://eduassistpro.github.io>

- 3 When interrupts happens, the interrupt p whatever it has to in order to handle the interrupt

- 4 Then it will unblock the driver that started it

Device Drivers

Device-specific code for controlling an I/O device

A driver for a mouse differs from a driver for a HDD

Handles one device type, or at most, one class of closely related devices

Part of

Positioning
control

Most OSs define a standard interface (between OS
for ~~block devices~~ and ~~character devices~~)

Must be flexible and be able to handle errors, several interfaces
etc.

Allowed to call only a handful of system calls, e.g. to allocate
memory for a buffer

Some parts are device-specific but others are device independent

Assignment Project Exam Help

There is no strict boundary between device-specific and device independent software and varies between OSs

Mos

- <https://eduassistpro.github.io>
- Buffering
- Error reporting
- Allocating and releasing dedicated device
- Providing a device-independent block size

Add WeChat edu_assist_pr

Assignment Project Exam Help

<https://eduassistpro.github.io>

Add WeChat edu_assist_pr

- (a) Without a standard driver interface
- (b) With a standard driver interface

Uniform Interfacing for Device Drivers II

Interface between the driver and OS is defined

OS can install new driver easily and the writer of the driver knows what it can expect from the OS

In practice, not all devices are absolutely identical, but there are only a small number of device types

For ea

funct

Ofte

functions

OS records the address of the table when the driver is loaded
indirect calls via this table

Another aspect of having a uniform interface is how I/O devices are named: each device has a **major device number** and **minor device number**

Closely related to naming is protection → devices appear as files in the file system, so usual protection rules could be used

Assignment Project Exam Help

<https://eduassistpro.github.io>

Add WeChat edu_assist_pro

- (a) Unbuffered input
- (b) Buffering in user space
- (c) Buffering in the kernel followed by copying to user space
- (d) Double buffering in the kernel

Ways to handle data streams from I/O devices

- Each interrupt may wake-up user's process
- OS writes into user space's buffer and wakes up user's process once the buffer is full. What happens if the buffer is paged

- <https://eduassistpro.github.io>

arrives at the time when the buffer is being copied to user's space?

- Double buffering
 - Circular buffering
- Buffering is also important for output → e.g. when sending data over a slow telephone line

Assignment Project Exam Help

<https://eduassistpro.github.io>

Add WeChat edu_assist_pr

Networking may involve multiple copies of a packet

Errors are far more common in the context of I/O than in other contexts

Many errors are device-specific and must be handled by appropriate driver, but the framework for error handling is device independent

Clas

- <https://eduassistpro.github.io>

there's only two

- Solution → just report back an error code
- Actual I/O errors → write to a disk block that is damaged or read from a camera that is turned off
 - Solution → it is up to the driver to decide what to do, whether to try to solve the problem or report back the error code

Assignment Project Exam Help

Different disks may have different sector sizes

It is up to

provi

<https://eduassistpro.github.io>

Some devices deliver data one byte at a time (e.g. most
others deliver theirs in larger units (e.g. network in

Add WeChat edu_assist_pr

Most of the I/O software is within the OS

Small portion of it consists of libraries linked together with user programs, and even whole programs running outside the kernel

Exa

Proc
call

Other procedures may do actual work: e.g. format

Not all user-level I/O software consists of library procedures → another important category is the spooling system

Some devices, such as CD-ROM recorders, can be used only by a single process at any given moment and cannot be shared

- Blocking user access to these devices causes delays and

bottlenecks

Spool to intermediate medium (disk file)

Spoo

mult

<https://eduassistpro.github.io>

Spooled devices (e.g. printers)

- 1 Printer output saved to disk file
- 2 File printed later by spooler daemon
 - Printer only allocated to spooler daemon
 - No normal process allowed direct access

- Provides sharing of non-sharable devices
- Reduces I/O time → gives greater throughput

Assignment Project Exam Help

<https://eduassistpro.github.io>

Add WeChat edu_assist_pr

Layers of the I/O system and the main functions of each layer

Loadable kernel modules provide device drivers

- Contain object code, loaded on demand
- Dynamically linked to running kernel

• <https://eduassistpro.github.io>

Kmod Add WeChat edu_assist_pr

- Kernel subsystem managing modules wi
- Determines module dependencies
- Loads modules on demand

Every LKM consists of two basic functions (minimum)

```
/* used for all initialisation code */  
int init_module (void) {  
    ...  
}  
/* use  
void cleanup_module (void) {  
    ...  
}
```

Load module by using the **insmod** command → normally restricted to root

Kernel provides common interface for I/O system calls

Assignment Project Exam Help

Devices grouped into device classes

- Members of each device class perform similar functions



<https://eduassistpro.github.io>

Major and minor identification numbers

- Used by device drivers to identify their device
- Devices with same major number controll

- Minor numbers enable system to distinguish between devices of same class

Most devices represented by device special files

Device files accessed via virtual file system (VFS) (`/dev`)

- System calls pass to VFS, which in turn issues calls to device drivers → most drivers implement common file operations,

List of

Linu

→ retrieving status information from printer

<https://eduassistpro.github.io>

Add WeChat edu_assist_pr

Character Device

Assignment Project Exam Help

- Transmits data as stream of bytes
- Represented by `device_struct` structure, which contains

- <https://eduassistpro.github.io>

`file_operations` structure

- Maintains operations supported by device
- Stores functions called by VFS when system call accesses device special file

Assignment Project Exam Help

<https://eduassistpro.github.io>

Add WeChat edu_assist_pr

Block I/O subsystem Assignment Project Exam Help

- Kernel's block I/O subsystem contains number of layers
-

<https://eduassistpro.github.io>

Two primary strategies used by kernel to minimise amount of time spent accessing block devices

- Caching data
- Clustering I/O operations

When data from block device requested, kernel first searches cache

- If found, data copied to process address space



Direct <https://eduassistpro.github.io>

- Driver bypasses kernel cache when access
- Important for databases and other applications
caching inappropriate and may reduce performance/consistency

Assignment Project Exam Help

<https://eduassistpro.github.io>

Socket (message)

Network interface

Add WeChat edu_assist_pr

Assignment Project Exam Help

```
fd = open (filename, permission)
```

- Opens file for reading/writing



- <https://eduassistpro.github.io>

```
fd = open (filename, mode)
```

- Mode is 0, 1, 2 for read, write, read/write

```
close (fd)
```

Add WeChat edu_assist_pro

Assignment Project Exam Help

```
numbytesread = read (fd, buffer, numbytes)
```

- Read `numbytes` from file or device referenced by `fd` into

- <https://eduassistpro.github.io>

```
numbyteswritten = write (fd, buffer, num
```

- Write `numbytes` to file referenced by `fd`
- Returns number of bytes actually written in `fd`

Add WeChat [edu_assist_pro](#)

```
pipe (&fd[2])
```

- Creates pipe

- `fd` is an array of 2 integers (`fd[0]` for reading, `fd[1]` for writing)

```
newf
```

- <https://eduassistpro.github.io>

```
ioctl (fd, operation, &termios)
```

- Used to control devices: e.g. file chars

```
fd = mknod (filename, permission, dev)
```

- Creates new special file e.g. character or block device

Each process has its own file descriptor table

- Each process has 3 file descriptors when created

1	
2	

By default, all three file descriptors refer to terminal from which program was started

Blocking I/O

Assignment Project Exam Help

- Call returns when operation completed
- Process suspended I/O appears “instantaneous”
-

<https://eduassistpro.github.io>

Non

- I/O call returns as much as available (e.g. read)
- Turn on for file descriptor using `fchmod` system
- Provides application-level polling for I/O

Add WeChat [edu_assist_pro](#)

Asynchronous I/O

Process executes in parallel with I/O operation

- No blocking in interface procedure

I/O sub-
com

- <https://eduassistpro.github.io>

...

Supports check/wait if I/O operation completed

Very flexible and efficient

Harder to use and potentially less secure